

24-HOUR LLM BUILD SPRINT

Hack Armenia

Build an LLM system that can **evaluate** evidence, **reason** under uncertainty, and **generalize** — for Armenia's public interest.

August 8–9, 2026 · AI9 Startup Campus · Yerevan

Welcome to Hack Armenia

One weekend. One theme. A working LLM system that actually does something useful for Armenia — and evidence that it works.



Build

24 hours to go from idea to a running prototype, with mentors in the room.



Measure

Show numbers. An evaluation — even a lightweight one — is what separates the top teams.



Pitch

5 minutes in front of an expert jury on Sunday. Talk plus live demo.

The sprint in numbers

24h
of building
Sat 10:00 → Sun

60–80
participants
one cohort

15–20
teams
3–4 people each

12
mentors
on-site Saturday

9
jury members
judging Demo Day

9
domains
Armenia's public interest

The theme: Armenia's public interest

Nothing is revealed at kickoff — the theme *is* the brief. Pick a real problem that matters here, and build for it.

Evaluate

Weigh evidence and cite what it is grounded in.

Reason

Work through uncertainty instead of guessing confidently.

Generalize

Hold up across real cases, not just your happy path.

The domains and example projects that follow are there to guide you — bring your own idea in the same spirit.

Nine public-interest domains

Where LLMs can be genuinely useful to Armenia. Choose one — or bring your own.



Urban development
City ops, planning & mobility



Public systems & policy
Filings, evidence & transparency



Healthcare
Safe triage, admin & patient education



Education
Tutoring, grading & Armenian-language learning



Security
Phishing defense, secure code & threat triage



Food & agriculture
Crop advisory, food safety & nutrition



Environment & energy
Efficiency, hazards & resources



Finance & fintech
Fair credit, fraud & literacy



Emergency & disaster
Coordination, preparedness & response

Example projects

Starting points, not requirements. Full list on the website.

🏛️

SME tax-filing copilot

Walk a small business through declarations; extract figures from receipts and pre-fill the forms.

🩺

Guarded symptom triage

Route symptoms to the right level of care in Armenian — with hard safety rails. Not diagnosis.

🎓

Auto-grading assistant

Grade short answers against a rubric, write feedback, and flag borderline cases for the teacher.

🔒

Phishing & scam detector

Catch bank-impersonation SMS and fake government notices aimed at Armenian users.

🍎

Crop advisory agent (VLM)

Diagnose plant disease from a photo and advise on treatment and timing, in Armenian.

🚨

Response coordination agent

Turn chaotic incoming incident reports into a de-duplicated, prioritised response picture.

Rules & tools

🧠

Any model or API

Open or closed, local or hosted. Bring your own keys where you can.

🔓

Open source required

Projects must be open source and fully reproducible. No closed black boxes.

🇦🇲

Armenian encouraged

Not required, but Armenian-language support is a real plus.

📦

Use real data

Anything publicly available, or data we provide before or during the event.

👥

Teams of 3–4

Form your team at kickoff. 15–20 teams in total.

⚖️

Build it here

Come with ideas, not with a finished codebase. The work happens during the sprint.

The 24 hours

Sat 10:00	**Kickoff & briefing** Welcome, theme walkthrough, logistics.
Sat 11:00	**Teams form & plan** Form a team of 3–4, scope your problem and data.
Sat 15:00–18:00	**Mentors on-site** Mentors available in the room — flexible, drop in any time.
Sat → Sun	**Build** Prototype end-to-end early, then harden and evaluate.
Sun 12:00	**Code freeze** Lock the repo, prepare your talk and demo.
Sun 13:00	**Demo Day** 5-minute pitches to the jury.
After demos	**Awards** The jury names the winners.

Your mentors

On-site **Saturday 15:00–18:00** — flexible, drop in any time.



Philipp Guevorguian
Technical Staff · Perceptron AI



Aram Shakhbandaryan
Founder · Empy.ai



Ani Vanyan
ML Researcher · YerevaNN



Khoren Petrosyan
ML Researcher · YerevaNN



Naré Gevorgyan
CEO · AICA



Marie Mikayelyan
CPO · OmniShift



Hovhannes Kuloghlyan
Co-Founder · Wirestock



Armen Grigoryan
Head of TV & Digital Products · Ucom



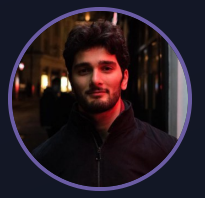
Khachik Badeyan
Technical Portfolio Lead · PMI Science Armenia



Gevorg Minasyan
Async



Davit Karamyan
NVIDIA



Tigran Atoyan
Founder & AI Researcher

How to use your mentors

Come with a specific question

“Is this evaluation sound?” beats “what do you think of our idea?”. Mentors are most useful once you have something concrete.

Show them something running

Even a broken prototype gives a mentor more to work with than a slide.

Ask about scope

The most common 24-hour failure is building too much. Mentors will help you cut.

Product and technical both

Some mentors will push on your model and evaluation, others on who actually uses this and why.

Mentors are on-site Saturday 15:00–18:00. Timing is flexible — grab whoever is free.

What you hand in

🎬

Working demo

It has to run. A rough end-to-end demo beats a polished deck.

📦

Artifact

A repo with reproducible code — plus an app link if you have one.

📊

Evaluation

A test set, rubric, examples, or benchmark — with numbers.

🎯

Impact note

Who uses it, and what changes if it works.

Code freeze is Sunday 12:00 — lock the repo and prepare your pitch before then.

The pitch: 5 minutes, maximum

Five minutes total per team, split however you like between the talk and the demo.

🎤 The talk

The problem, why an LLM is the right tool, how you evaluated it, and what you found — including where it breaks.

🎬 The demo

Show it running on a real case. Live is best; a recording is fine if the network is not.

Split it as you see fit — but at 5:00 you stop. Practise it once before you present.

How you are judged

Five criteria, one rubric for every team — whatever domain you pick.

1. **Idea quality and public-interest fit**
   Is this a real Armenian problem, with a real user and a benefit you can point at?

2. **Working prototype and appropriate LLM use**
   It has to run — and the model has to be doing real work, not decoration.

3. **Evaluation and supporting evidence**
   Numbers, not vibes. Show how you measured it and what came out.

4. **Model limitations, failure modes, and safety**
   Know where it breaks, who it could harm, and what you do about it.

5. **Real-world feasibility and reproducibility**
   Could this survive outside the demo — and can someone else run it?

0 — missing · **3 — adequate** · **5 — exceptional** · **25 points total**

The rubric at a glance

Criterion-specific evidence anchors.

Criterion	0 points	3 points — adequate	5 points — exceptional
1. Idea quality and public-interest fit	Problem is unclear, outside scope, or has no identified user.	A clear Armenia-relevant problem and user are identified; the benefit is plausible.	Local need and user workflow are well understood, with specific and measurable potential impact.
2. Working prototype and appropriate LLM use	No working prototype is shown, or the LLM/VLM is merely decorative.	The core workflow works on a representative case, and the model's role is credible.	The system works end to end across varied cases and shows why an LLM/VLM is better than a simpler alternative.
3. Evaluation and supporting evidence	No meaningful evaluation evidence is provided.	An appropriate test set, examples, metric, or rubric is used and results are reported.	Evaluation is realistic and reproducible, includes a baseline or error analysis, and supports the team's claims.
4. Model limitations, failure modes, and safety	Relevant limitations, failure modes, or safety risks are not identified.	Important risks and failure cases are identified, with basic safeguards or human escalation.	Risky cases are tested systematically and effective mitigation, abstention, escalation, or monitoring is demonstrated.
5. Real-world feasibility and reproducibility	The project relies on toy assumptions, is not reproducible, or has no credible path beyond the demo.	It uses realistic data or labelled stand-ins, documents setup, and describes a plausible deployment path.	It fits real workflows and considers integration, cost, latency, privacy, and maintenance; reproduction is straightforward.

1 · Idea quality and public-interest fit

Is this a real Armenian problem, with a real user and a benefit you can point at?

0 POINTS

Problem is unclear, outside scope, or has no identified user.

3 — ADEQUATE

A clear Armenia-relevant problem and user are identified; the benefit is plausible.

5 — EXCEPTIONAL

Local need and user workflow are well understood, with specific and measurable potential impact.

WHAT THIS LOOKS LIKE

A construction-permit pre-check. **3** — names municipal reviewers as the user and explains the pain. **5** — shows how many applications bounce today, which rejection reasons dominate, and what share the tool would catch before filing.

2 · Working prototype and appropriate LLM use

It has to run — and the model has to be doing real work, not decoration.

0 POINTS

No working prototype is shown, or the LLM/VLM is merely decorative.

3 — ADEQUATE

The core workflow works on a representative case, and the model's role is credible.

5 — EXCEPTIONAL

The system works end to end across varied cases and shows why an LLM/VLM is better than a simpler alternative.

WHAT THIS LOOKS LIKE

A symptom-triage assistant. **3** — runs one vignette end to end. **5** — runs ten varied vignettes and shows a keyword baseline mis-routing the cases the model gets right.

3 · Evaluation and supporting evidence

Numbers, not vibes. Show how you measured it and what came out.

0 POINTS

No meaningful evaluation evidence is provided.

3 — ADEQUATE

An appropriate test set, examples, metric, or rubric is used and results are reported.

5 — EXCEPTIONAL

Evaluation is realistic and reproducible, includes a baseline or error analysis, and supports the team's claims.

WHAT THIS LOOKS LIKE

An auto-grading assistant. **3** — reports agreement with teacher marks on 30 labelled answers. **5** — adds a majority-class baseline, breaks errors down by question type, and ships a script that re-runs the whole thing.

4 · Model limitations, failure modes, and safety

Know where it breaks, who it could harm, and what you do about it.

0 POINTS

Relevant limitations, failure modes, or safety risks are not identified.

3 — ADEQUATE

Important risks and failure cases are identified, with basic safeguards or human escalation.

5 — EXCEPTIONAL

Risky cases are tested systematically and effective mitigation, abstention, escalation, or monitoring is demonstrated.

WHAT THIS LOOKS LIKE

The same triage assistant. **3** — names under-triage as the key risk and always offers a “seek care” fallback. **5** — measures the under-triage rate on deliberately hard vignettes and abstains instead of guessing when confidence is low.

5 · Real-world feasibility and reproducibility

Could this survive outside the demo — and can someone else run it?

0 POINTS

The project relies on toy assumptions, is not reproducible, or has no credible path beyond the demo.

3 — ADEQUATE

It uses realistic data or labelled stand-ins, documents setup, and describes a plausible deployment path.

5 — EXCEPTIONAL

It fits real workflows and considers integration, cost, latency, privacy, and maintenance; reproduction is straightforward.

WHAT THIS LOOKS LIKE

An e-invoice extractor. **3** — real invoice layouts and a README that actually runs. **5** — reports cost and latency per document, says where personal data is handled, and names the system it would plug into.

The jury

Who you present to on Sunday at 13:00.



Gevorg Mantashyan

First Deputy Minister · Ministry of High-Tech Industry



David Zokhrabyan

Growth · OpenAI



Davit Arakelyan

Chief Business Officer · Eleveight AI



Arik Akhverdyan

Co-Founder & CEO · AI9



Vazgen Hakobjanyan

Founder & CEO · Magical Labs



Vazgen Badalyan

Founder · Evidenso



Erik Arakelyan

Senior Researcher · NVIDIA Armenia



Marie Mikayelyan

CPO · OmniShift



Arto Yeritsyan

CEO & Founder · Async

The organizers

Hack Armenia is run by the Armenia LLM Summer School, AI9 and YerevaNN.



Armen Aghajanyan
Perceptron AI



Aram Galstyan
USC / Amazon



Erik Arakelyan
Senior Researcher · NVIDIA Armenia



Hrant Khachatryan
Founding Director · YerevaNN



Perouz Taslakian
ServiceNow AI Research



Tatev Vardanyan
YerevaNN



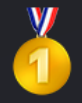
Arik Akhverdyan
Co-Founder & CEO · AI9



Roza Khachatryan
COO · AI9 Startup Campus

Prizes

OpenAI is sponsoring the hackathon prizes with API credits.



1st place

$12,000

in OpenAI API credits



2nd place

$7,500

in OpenAI API credits



3rd place

$5,000

in OpenAI API credits

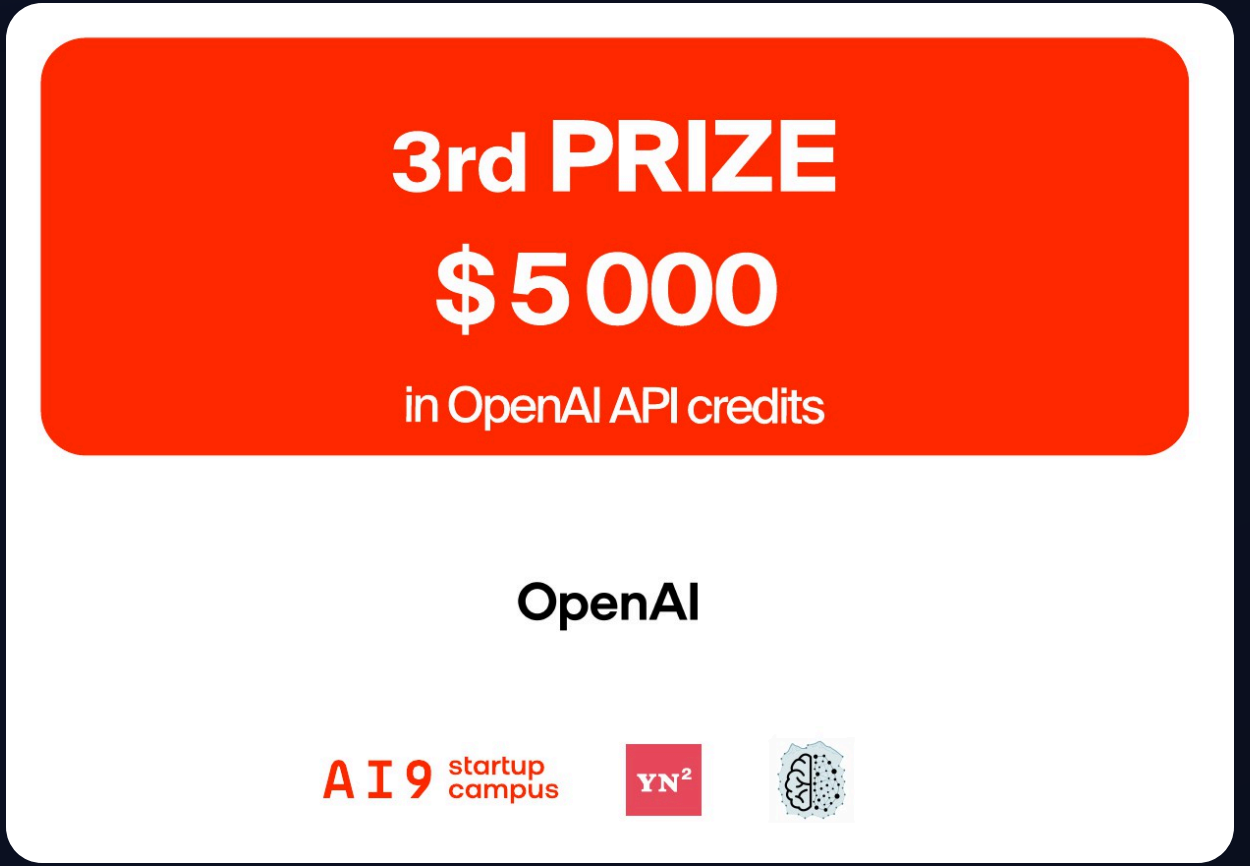
1st prize


1st PRIZE
$12 000
in OpenAI API credits
OpenAI
AI9 startup campus
YN²

2nd prize


2nd PRIZE
$7500
in OpenAI API credits
OpenAI
AI9 startup campus
YN²

3rd prize


3rd PRIZE
$5 000
in OpenAI API credits
OpenAI
AI9 startup campus
YN²

Supported by

🏛️

Ministry of High-Tech Industry

Hack Armenia is generously supported by the Ministry of High-Tech Industry of the Republic of Armenia.

🤝

ISAA

Supporting the sprint alongside the Ministry.

Special prize

A separate award, presented by Eleveight AI.



$5,000 equivalent GPUs

Eleveight AI is presenting its own special prize at Hack Armenia — GPUs worth $5,000, so the winning team leaves with real compute to keep building.

NOW GO BUILD

Good luck

Ship something that runs, measure it honestly, and tell us who it helps.

Demo Day — Sunday 13:00 · armeniallm@gmail.com

